Scientific Review: Recursive Deep Models for Semantic Compositionality Over a Sentiment Treebank

Mounika Marreddy
University of Bonn, Germany
mounika0559@gmail.com
Github Link: https://github.com/mounikamarreddy/UKP-Assignment

Background & Motivation. Given two sentences with the same words but different word order: [+ white blood cells destroying an infection, - an infection destroying white blood cells] [1]. Traditional bag-of-words or TF-IDF models struggle to distinguish the sentiment between these sentences. These models do not capture word order or the relational meaning between words, which is crucial for understanding the sentiment conveyed by the sentence structure. Additionally, these studies categorize the longer documents as positive or negative based on the presence of more positive or negative words. However, identifying sentiment in short phrases with these simple word features presents several challenges: (1) Short sentences often lack sufficient context, (2) These methods are unable to handle words with different meanings, and (3) They fail to capture the tone in short sentences.

The current study draws inspiration from cognitive linguistic studies to better understand the meaning of short phrases or sentences. One particularly influential study that I noticed in the field of cognitive neuroscience proposed by Pallier et al. (2011), found increased brain activity when adding a word to a constituent to form a meaningful sentence. This study demonstrates that the brain processes hierarchical structures and word meanings. Similarly, this work also focuses on capturing the meaning of short phrases to improve sentiment accuracy at both the sentence and document levels.

Challenges. 1. One major limitation in performing sentiment analysis at the short text or sentence level is the unavailability of annotated datasets at the phrase level. To address this issue, the authors created an annotated dataset with phrase-level sentiment labels. Specifically, the authors focused on creating labels that are not just binary but more fine-grained. 2. Current neural network models, like Recurrent Neural Networks (RNNs) and Recursive Neural Networks (RecNNs), have limitations in efficiently handling sequential and hierarchical data. To overcome these, the authors introduced the Recursive Neural Tensor Network (RNTN), which uses a single composition function with shared weight parameters applied recursively. This model captures the relationships between words and phrases more effectively, making it ideal for tasks like sentiment analysis, where syntactic grouping and word order are crucial. Overall, the RNTN structure resembles RecNNs, but its complexity depends on how the words in a sentence are formed into a parse tree within the network.

Key Contributions of the Paper

- The Stanford Sentiment Treebank (SST), a large dataset with phrase-level annotations, was introduced to capture sentiment in short texts and enable fine-grained sentiment analysis.
- Introduced RNTN model, designed based on the parse tree structure of sentences, which allows for a more comprehensive understanding of the relationships between words within a sentence.
- Experiment with several previous approaches, including RecNN and MV-RecNN, as well as the proposed RNTN method. Provide benchmark performance on the SST dataset for binary sentiment classification and fine-grained sentiment analysis.

Discussion

Pros

- The proposed sentiment tree bank dataset offers several benefits: (1) Deeper understanding of sentiment allows for fine-grained aspects, capturing sentiments from very positive to very negative. (2) The phrase-level annotations allow us to provide deeper insights into sentiment across different parts of the text. (3) The syntactic parse trees help understand the syntactic structure of sentences.

[1]https://cs224d.stanford.edu/lectures/CS224d-Lecture11.pdf

- The proposed RNTN model seems appropriate and sound. In particular, the model learns both the word and phrase representations recursively and predicts sentiment labels at every phrase, resulting in better accuracy.

Limitations

- One limitation of the current approach is that the author's motivation towards only the sentiment analysis task. However, how one can use the semantic word or phrase representations can be used in complex natural language processing tasks like machine translation, question answering where word order and syntactic structure and multi-word expressions are very important.
- Second, the current SST dataset is created only for English. However, the phrase structure for other languages can be different due to morphological variations. Replicating similar datasets for other languages requires linguistic expertise to annotate phrase structures accurately and ensure they capture the nuances of those languages. Additionally, while the RNTN model may work for different languages, it would need to be adapted to handle the language's specific syntactic and morphological characteristics. This could involve modifying the model to better account for language-specific features and training it on annotated datasets particular to those languages.
- Third, it is important to note that while the current work focused on the compositional and hierarchical structure of sentences using a newly created dataset, there is a question of whether the RNTN model captures long-term dependencies effectively. The complexity of compositional structures increases with longer texts, such as books or narrative stories, where handling these requires observing long-term semantic context. Ensuring the model can manage these extended dependencies is crucial for accurately processing and understanding longer, more complex texts.

Future Improvements

- The proposed SST dataset is useful for further understanding emotion, hate-speech, and discourse. Since the current study focuses only on sentiment, it would be beneficial to transfer this knowledge to other existing sentiment datasets at fine-grained and binary class levels. Additionally, creating a single model for all tasks, rather than a complex model for each specific task, would make it more versatile and applicable to a wider range of downstream NLP tasks.
- Interpretation of the word and phrase structure embeddings can be achieved through probing tasks, similar to recent studies proposed in Tenney et al. (2018); Jawahar et al. (2019). These tasks help to identify the linguistic properties encoded in the representations, such as syntactic roles, semantic meanings, and hierarchical relationships within the language. By analyzing these embeddings, we can determine if the model captures language hierarchy, including the relationships between different levels of linguistic structure, such as phrases, clauses, and sentences. This analysis can reveal the extent to which the model understands and represents the intricate details of the language.
- Word and phrase structure embeddings from the RNTN model can provide insights into brain function in the following ways: (1) Hierarchical Syntactic Processing: By utilizing naturalistic story brain datasets, we can analyze how the brain processes hierarchical syntactic structures and identify the brain regions involved in word-level, phrase-level, and sentence processing. (2) Syntactic vs. Semantic Information: The RNTN model can disentangle syntactic and semantic information, allowing us to study how the brain processes these different types of information and determine which regions are responsible for each (Reddy and Wehbe, 2021).

References

Ganesh Jawahar, Benoît Sagot, and Djamé Seddah. 2019. What does bert learn about the structure of language? In *ACL*.

Christophe Pallier, Anne-Dominique Devauchelle, and Stanislas Dehaene. 2011. Cortical representation of the constituent structure of sentences. *PNAS*.

Aniketh Janardhan Reddy and Leila Wehbe. 2021. Can fmri reveal the representation of syntactic structure in the brain? *NeurIPS*, 34:9843–9856.

Ian Tenney, Patrick Xia, et al. 2018. What do you learn from context? probing for sentence structure in contextualized word representations. In *ICLR*.